

# DynamicEarthNet: towards FedSPD

April 1, 2025

# Overview

Data setting

Motivation for Federated Learning

Federated Procedure

- Riemannian local optimization

- Differential Privacy Option

- Euclidean local optimization

Server Aggregation

- Riemannian aggregation

- Euclidean aggregation

Average gradient stream approach

# DynamicEarthNet dataset

See [TKW<sup>+</sup>22].








Multi-spectral images

- ▶ 75 Areas of Interest (AOI)
- ▶ Satellite images collection: daily 2018/01/01  $\rightarrow$  2019/12/31
- ▶ 730 images/AOI  $\rightarrow$  54,750 satellite images in total
- ▶ 4 channels: RGB+NIR
- ▶ Spatial resolution: 3 meters
- ▶ Image size:  $1024 \times 1024$

One time-series:  $x \in \mathbb{R}^{T \times H \times W \times 4}$  with  $T$ : number of days in the month,  
 $H = W = 1024$ .

# Land Use/Land Cover (LULC)

- ▶ Labels:  $24 \times 75 = 1,800$
- ▶ 7 different (imbalanced) classes:

class name	%	#AOIs	color
impervious surface	7.1	70	
agriculture	10.3	37	
forest & other vegetation	44.9	71	
wetlands	0.7	24	
soil	28.0	75	
water	8.0	58	
snow & ice	1.0	2	

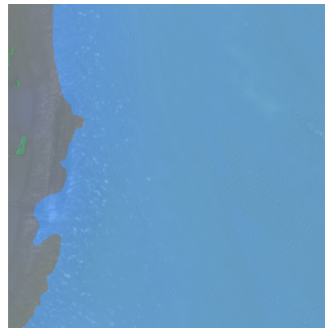
No single AOI contains all 7 classes.

For a given time series, label:

$$y \in \{0, \dots, 6\}^{T \times H \times W}$$



(a)



(b)

Figure: Two random images from different time series id (and different directories)

## Challenges:

Issue 1 Voluminous data (hundreds of GB)

Issue 2 Slow to move all images to a single processing server

# Why Federated Learning is relevant?

**Solution Issue 1** client  $\rightarrow$  smaller dataset & own stat. heterogeneity with clear interpretation

**Solution Issue 2** No exchange of raw datasets, just each client's updates

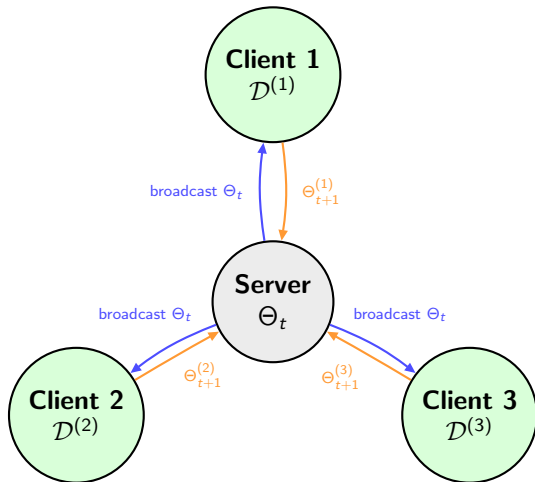
- ▶ Each AOI  $\rightarrow$  local phenomena
- ▶ FL: adapt to local distrib and aggregate knowledge for better generalization

★ **Idea 1** one AOI = one client  $\rightarrow$  too many clients (75);

★ **Idea 2** create clients based on dominant LULC categories across AOI

1. Urban areas
2. Agricultural areas
3. Forest-dominated areas
4. Water-dominated areas
5. Mixed land-use areas

# Federated Learning Overview



$K$  clients. Each client  $k$  has dataset  $D^{(k)}$  of  $n_k$  pixel annotation pairs  $\{\mathbf{X}_i^{(k)}, y_i^{(k)}\}_{i=1, \dots, n_k}$ , where

- ▶  $\mathbf{X}_i^{(k)}$  is a covariance matrix in  $\mathbb{R}^{4T \times 4T}$  (with  $T = 31$ ),
- ▶  $y_i^{(k)} \in \{0, \dots, 6\}$  is a semantic mask


Total number of samples  $n := \sum_{k=1}^K n_k$

Loss function  $\ell$  (CE)



Forward pass: SPD-Net  $f_{\Theta}(\cdot)$  (see [HVG17, BSB<sup>+</sup>19])

### SPD Feature Extraction



Each client's SPD-net with  $L$  layers has learnable parameters

$$\Theta = (\underbrace{W_1, \dots, W_L}_{\text{BiMap}}, \underbrace{G_1, \dots, G_L}_{\text{BN}}, \underbrace{\xi_1, \dots, \xi_d}_{\text{FC}}) \in \mathcal{M}$$

with

$$\mathcal{M} := \left( \prod_{l=1}^L \text{St}(d_l, d_{l+1}) \right) \times \left( \prod_{l=1}^L \text{SPD}_{d_l} \right) \times \mathbb{R}^d.$$

Local objective functions For each client  $k = 1, \dots, K$ ,

$$F_k(\Theta) := \frac{1}{n_k} \sum_{i=1}^{n_k} \ell(f_{\Theta}(x_i^{(k)}), y_i^{(k)}).$$

Global objective function

$$\arg \min_{\Theta \in \mathcal{M}} F(\Theta) := \sum_{k=1}^K p_k F_k(\Theta), \quad p_k = \frac{n_k}{n}.$$

# Federated procedure

See [LM22].

## Algorithm

- ★ **Initialization**  $\Theta_0$  to all clients (random)
- ★ **Outer loop** For  $t = 0, \dots, T - 1$  (rounds):
  1. A subset  $\mathcal{S}_t$  of  $r \leq K$  clients is sampled (random)
  2. Server broadcasts  $\Theta_t$  to clients.
  3. Each client  $k \in \mathcal{S}_t$  runs  $E_k$  local updates on  $\Theta$  with initialization  $\Theta_t$ .
  4. Aggregation of each client's local updated parameters via barycenters  $\rightarrow \Theta_{t+1}$
- ★ **Global output** Final aggregated param.  $\Theta_T$  or  $\text{Unif}(\{\Theta_0, \dots, \Theta_T\})$

If each local function  $F_k$  is  $L_k$ -smooth and geodesically convex, then the algorithm converges in  $O(1/T)$  (provided some curvature/manifold additional assumptions).

# Riemannian local updates at client $k \in \mathcal{S}_r$

Let  $\theta$  denote a Riemannian parameter (either one of the  $W_l$  or  $G_l$ ).

- ▶ Initialize  $\theta_0^{(k)} = \theta_t$
- ▶ For  $s = 1, \dots, E_k$ :

$$\theta_{s+1}^{(k)} = R_{\theta_s^{(k)}} \left( -\eta^{(k)} g^{(k)} \right),$$

where the local Riemannian gradient is

$$g^{(k)} = \text{grad } F_k(\theta_s^{(k)}) - \mathcal{T}_{\theta_t \rightarrow \theta_s^{(k)}} \left( \text{grad } F_k(\theta_t) - \text{grad } F(\theta_t) \right),$$

- ▶ The blue term is the optional **SVR correction** (variance reduction).

# Differential Privacy as a model variant

Besides the [variance-reduction](#) approach, we can also inject **DP noise** (see [HHJM24a, HMJG24]) on each local gradient to protect sensitive data.

**Key idea:** Add tangent-space Gaussian noise before retraction/exponential update.

- ▶ **DP** can be combined with or without the [SVR](#) term.
- ▶ Each client  $k$  ensures local DP by perturbing its gradient:

$$g_{\text{DP}}^{(k)} = \text{clip}_{\tau}(g^{(k)}) + \varepsilon_t, \quad \text{where } \varepsilon_t \sim \mathcal{N}_{\theta_s^{(k)}}(0, \sigma^2).$$

$\mathcal{N}_{\theta_s^{(k)}}(0, \sigma^2)$  is *tangent-space Gaussian* distribution at  $\theta_s^{(k)}$  and  $\text{clip}_{\tau} : T_{\theta}\mathcal{M} \rightarrow T_{\theta}\mathcal{M}$  is such that

$$\text{clip}_{\tau}(\xi) = \min \left\{ \frac{\tau}{\|\xi\|_{\theta}}, 1 \right\}.$$

**Privacy Guarantee:** With an appropriate choice of  $\sigma$  and composition rules, the entire FL pipeline can provide an  $(\varepsilon, \delta)$ -DP guarantee.

# Riemannian local updates with optional DP

- ▶ Initialize  $\theta_0^{(k)} = \theta_t$
- ▶ For  $s = 1, \dots, E_k$ : over a selected minibatch,

$$g^{(k)} = \text{grad } F_k(\theta_s^{(k)}) - \mathcal{T}_{\theta_t \rightarrow \theta_s^{(k)}} \left( \text{grad } F_k(\theta_t) - \text{grad } F(\theta_t) \right),$$

$$g_{\text{DP}}^{(k)} = \text{clip}_\tau(g^{(k)}) + \varepsilon_s, \quad \varepsilon_s \sim \mathcal{N}_{\theta_s^{(k)}}(0, \sigma^2).$$

Then update via retraction (or exponential map):

$$\theta_{s+1}^{(k)} = R_{\theta_s^{(k)}} \left( -\eta^{(k)} g_{\text{DP}}^{(k)} \right).$$

- ▶ **Red** term is the **DP noise** addition (optional)
- ▶ **Blue** term is still the **SVR correction** (optional).

# Local updates for Euclidean FC weights (SVR + DP)

Let  $\xi \in \mathbb{R}^{d_{\text{FC}}}$  be the parameter vector for the final fully-connected layer. During each local epoch on client  $k$ :

$$\xi_0^{(k)} = \xi_t, \quad \text{and for } s = 0, \dots, E_k - 1 :$$

$$\xi_{s+1}^{(k)} = \xi_s^{(k)} - \eta^{(k)} \left( \nabla F_k(\xi_s^{(k)}) - [\nabla F_k(\xi_t) - \nabla F(\xi_t)] + \varepsilon_s \right),$$

where:

- ▶  $\nabla F_k(\xi_s^{(k)})$ : local Euclidean gradient on client  $k$ .
- ▶  $[\nabla F_k(\xi_t) - \nabla F(\xi_t)]$  is the **SVR correction** term.
- ▶  $\varepsilon_s \sim \mathcal{N}(0, \sigma^2 I)$  is the **DP noise** term.
- ▶ After  $E_k$  steps, client  $k$  sends  $\xi_{E_k}^{(k)}$  to the server.

# Server aggregation (Karcher flow) for $\theta = G$ SPD

After each client  $k \in \mathcal{S}_r$  sends back  $\theta_{E_k}^{(k)}$ , we compute the Riemannian (Fréchet) mean on the manifold:

$$\theta_{t+1} = \arg \min_{\theta \in \mathcal{M}} \sum_{k \in \mathcal{S}_r} p_k d(\theta, \theta_{E_k}^{(k)})^2.$$

► In practice, we use a gradient-descent-based procedure (Karcher flow):

$$\theta_{t+1} \leftarrow \text{Exp}_{\theta_t} \left( -\gamma \sum_{k \in \mathcal{S}_r} p_k \text{Log}_{\theta_t}(\theta_{E_k}^{(k)}) \right),$$

where  $\gamma$  is a step size for the Karcher flow.



# Server aggregation (projection-based) on $\mathbf{W} \in \text{St}(d_I, d_{I+1})$

See [BLSB25].

For the BiMap parameters  $W_I$  living on the Stiefel manifold, we can use:

- **Karcher flow** or Riemannian gradient-based aggregator:

$$W_{t+1} = \arg \min_{W \in \text{St}(d_I, d_{I+1})} \sum_{k \in S_t} p_k d^2(W, W_\tau^{(k)}),$$

which can be iteratively approximated (exp/log maps).

- **Projection-based aggregator:**

$$W_{t+1} = \underbrace{\text{uf}\left(\sum_{k \in S_t} p_k W_\tau^{(k)}\right)}_{\text{arithmetic mean}} \quad \text{or} \quad W_{t+1} = \text{qf}\left(\sum_{k \in S_t} p_k W_\tau^{(k)}\right).$$

No iterative scheme needed, computationally cheaper (simple retractions/liftings).

# Server aggregation of FC weights

At round  $t$ , we receive each client's updated  $\xi_{E_k}^{(k)}$  for  $k \in \mathcal{S}_t$ . To get  $\xi_{t+1}$ , standard weighted average (FedAvg):

$$\xi_{t+1} = \sum_{k \in \mathcal{S}_t} \frac{n_k}{\sum_{j \in \mathcal{S}_t} n_j} \xi_{E_k}^{(k)}$$

No need for any Riemannian retraction or manifold-based aggregator.

# Average gradient stream approach

See [HHJM24b].

Key ideas :

1. Aggregate each client's mini-batch gradient with vector transport to server.
2. Update global model.

Pros :

- ▶ No communication of parameters  $\rightarrow$  no expensive exp/log maps.
- ▶ Works for general manifolds (not just embedded or compact manifolds).
- ▶ Sublinear convergence for non-convex objectives (fixed step-size).

Cons :

- ▶ High variance of aggregating gradients, cost of vector transport
- ▶ Memory overhead on clients (store and transmit cumulative transported gradients)

# RFedAGS algorithm

★ **Outer loop** For  $t = 0, \dots, T - 1$ , server broadcasts  $x_t$ .

Each client  $k = 1, \dots, K$  does  $\theta_0^{(k)} \leftarrow x_t$  ;  $g_0^{(k)} \leftarrow 0$ .

1. For  $s = 0, \dots, E_k$ :

- Sample mini-batch  $B_s^{(k)}$ , compute local gradient

$$\eta_s^{(j)} \leftarrow -\alpha_t^{(k)} \frac{1}{\#B_s^{(k)}} \sum_{s \in B_s^{(k)}} \text{clip}_{\tau} \text{grad} f(\theta_s^{(k)}) + \varepsilon_s^{(k)}$$

2. Retraction-based update:  $\theta_{s+1}^{(k)} = R_{\theta_s^{(k)}}(\eta_{t,s}^{(k)})$

3. Transport gradient to server's tangent space and accumulate:

$$g_{s+1}^{(k)} = g_s^{(k)} + \mathcal{T}_{x_{t,s}^{(k)} \rightarrow x_t}(\eta_s^{(k)}).$$

★ **Aggregation** : Each client uploads  $g_{E_k}^{(k)}$  and server update:

$$x_{t+1} \leftarrow R_{x_t} \left( \frac{1}{K} \sum_{k=1}^K g_{E_k}^{(k)} \right).$$

★ **Output**  $x_T$ . In *red*, optional DP-RSGD style.

# References



Florent Bouchard, Nils Laurent, Salem Said, and Nicolas Le Bihan.

Beyond r-barycenters: an effective averaging method on stiefel and grassmann manifolds.  
*arXiv preprint arXiv:2501.11555*, 2025.



Daniel Brooks, Olivier Schwander, Frédéric Barbaresco, Jean-Yves Schneider, and Matthieu Cord.

Riemannian batch normalization for spd neural networks.  
*Advances in Neural Information Processing Systems*, 32, 2019.



Zhenwei Huang, Wen Huang, Pratik Jawanpuria, and Bamdev Mishra.

Federated learning on riemannian manifolds with differential privacy.  
*arXiv preprint arXiv:2404.10029*, 2024.



Zhenwei Huang, Wen Huang, Pratik Jawanpuria, and Bamdev Mishra.

Riemannian federated learning via averaging gradient stream.  
*arXiv preprint arXiv:2409.07223*, 2024.



Andi Han, Bamdev Mishra, Pratik Jawanpuria, and Junbin Gao.

Differentially private riemannian optimization.  
*Machine Learning*, 113(3):1133–1161, 2024.



Zhiwu Huang and Luc Van Gool.

A riemannian network for spd matrix learning.  
In *Proceedings of the AAAI conference on artificial intelligence*, volume 31, 2017.



Jiaxiang Li and Shiqian Ma.

Federated learning on riemannian manifolds.  
*arXiv preprint arXiv:2206.05668*, 2022.



Aysim Toker, Lukas Kondmann, Mark Weber, Marvin Eisenberger, Andrés Camero, Jingliang Hu, Ariadna Pregel Hoderlein, Çağlar Şenaras, Timothy Davis, Daniel Cremers, et al.

Dynamicearthnet: Daily multi-spectral satellite dataset for semantic change segmentation.